

# 一起动手学OR新特性



@Gerrard

# 内容

ngx.semaphore 信号量同步

ngx.shared.DICT 共享内存list

balancer\_by\_lua\* 用lua写负载均衡

stream-lua-nginx-module lua高性能tcp服务

resty 相关工具

实验环境 : nginx version: openresty/1.11.2.1

stream-lua-nginx-module: commit 75e0b16

相关代码 : <https://github.com/openresty-suzhou/meetup>

# 信号量ngx.semaphore

- 高效
- 进程内 跨context

文档 <https://github.com/openresty/lua-resty-core/blob/master/lib/nginx/semaphore.md>

# 生产者消费者问题

## 场景

- Coroutine A 生产数据 送入buffer
- Coroutine B 消费数据 从buffer取数据


## 原来的方式

Coroutine B 不断地polling + nginx.sleep

## 缺点

仍可能有毫秒级别的延迟

# 使用ngx.semaphore机制



```
local co = ngx.thread.spawn(produce, buffer)
local ok, err = sema:wait(1) --wait at most 1s
if ok then
    if buffer.message then
        ngx.say('main thread : receive message : ',
buffer.message)
    else
        ngx.say('main thread : unexpected error,
got semaphore but no message in buffer')
    end
else
    ngx.say('main thread : failed to wait on sema:
', err)
end
ngx.say('main thread : end.')
```

```
local function produce(buffer)
    ngx.say('producer thread : producing
data ...')
    ngx.sleep(0.1) --wait a little time
    buffer.message = 'This is a message
from producer!' --producing data
    ngx.say('producer thread : posting to
sema...')
    sema:post(1)
end
```

```
gerrard@ubuntu:~/workspace/test$ curl 127.0.0.1:8000/semaphore
main thread : no message in buffer, waiting for semaphore from producer...
producer thread : producing data ...
producer thread : posting to sema...
main thread : receive message : This is a message from producer!
main thread : end.
```

# ngx.shared.DICT 新增API

- lpush
- lpop
- rpush
- rpop
- llen

文档 <https://github.com/openresty/lua-nginx-module#ngxshareddict>

# 像redis的list ?


## 支持类型

- 字符串
- 数字

## 使用场景

- 轻松实现跨worker的stack, queue, deque

# 借助lpush,lpop封装一个简单的栈



```
function _M.new(self, name)
    self.name = name
    self.list = ngx.shared.list
    return setmetatable({}, {__index =
_M})
end
```


```
function _M.push(self,v)
    return self.list:lpush(self.name, v)
end
```

```
function _M.pop(self)
    return self.list:lpop(self.name)
end
```

```
function _M.size(self)
    return self.list:llen(self.name)
end
```



# 借助Ipush,Ipop封装一个简单的栈



```
local myStack = stack:new('stack')
for k, v in pairs(myStack) do
    print(k, tostring(v))
end

local elements = {'first', 2,true,4.001}

for _, v in ipairs(elements) do
    local len, err = myStack:push(v)
    if len then
        ngx.say('pushed value : ' .. v .. ' to stack,
stack length : ' .. len )
    else
        ngx.say('pushed to stack falied, err : ', err)
    end
end
```

```
while true do
    local val, err = myStack:pop()
    if not val then
        ngx.say('pop from stack falied, err : ', err)
        break
    end
    ngx.say('pop from stack, value : ' .. val .. ', sizeof
stack : ', myStack:size())
end
```

```
gerrard@ubuntu:~/workspace/test$ curl 127.0.0.1:8000/shdict
pushed value : first to stack, stack length : 1
pushed value : 2 to stack, stack length : 2
pushed to stack falied, err : bad value type
pushed value : 4.001 to stack, stack length : 3
pop from stack, value : 4.001, sizeof stack : 2
pop from stack, value : 2, sizeof stack : 1
pop from stack, value : first, sizeof stack : 0
pop from stack falied, err : nil
```

# 负载均衡 balancer\_by\_lua\*

- 动态选择upstream 细化到每个请求
- 灵活，自定义策略

文档 <https://github.com/openresty/lua-resty-core/blob/master/lib/nginx/balancer.md>

# upstream


## 原生upstream

- 手动配置，无法动态调整
- 策略有限 RR, weight, IP hash

## 使用场景

- 动态调整upstream
- 自定义负载均衡策略

# 实现一个简单的RR




```
local balancer = require 'ngx.balancer'
local host = '127.0.0.1'
local port = {9000, 9001, 9002}
local val, err = ngx.shared.list:incr('count', 1, -1)
if not val then
    ngx.log(ngx.ERR, 'failed to incr key "count", err : ', err)
    ngx.exit(500)
end

local state, status = balancer.get_last_failure()
if state then
    ngx.log(ngx.ERR, 'last peer failure:', state, " ", status)
end
```

```
if not ngx.ctx.tries then
    ngx.ctx.tries = 0
end
if ngx.ctx.tries < 1 then
    local ok, err = balancer.set_more_tries(1)
    if not ok then
        return error('failed to set more tries: ', err)
    elseif err then
        ngx.log(ngx.ERR, "set more tries:", err)
    end
end
ngx.ctx.tries = ngx.ctx.tries + 1
local ok, err = balancer.set_current_peer(host,
port[val%3+1])
if not ok then
    ngx.log(ngx.ERR, 'failed to set the current peer : ', err)
    return ngx.exit(500)
end
```

# 实现一个简单的RR



```
server {
    listen 8080;
    location / {
        proxy_pass http://backend/;
        proxy_next_upstream_tries 2;
    }
}
upstream backend{
    server 1.2.3.4;
    balancer_by_lua_file src/balancer.lua;
    keepalive 60;
}
```

```
server {
    listen 127.0.0.1:9000;
    location / {
        echo "upstream from port 9000...";
    }
}
server {
    listen 127.0.0.1:9001;
    location / {
        echo "upstream from port 9001...";
    }
}
server {
    listen 127.0.0.1:9002;
    location / {
        echo "upstream from port 9002...";
    }
}
```

# nginx stream

- 2015-04-28 nginx-1.9.0 mainline version has been released, with the stream module for **generic TCP/UDP proxying and load balancing**.

[http://nginx.org/en/docs/stream/nginx\\_stream\\_core\\_module.html](http://nginx.org/en/docs/stream/nginx_stream_core_module.html)

- 在这个特性之前nginx只能做7层的应用层代理, 4层的TCP代理一般用LVS, Haproxy等完成

# nginx stream

```
## stream的配置示例, 分别是udp proxy, tcp proxy, unix domain proxy
stream {
    upstream backend {
        hash $remote_addr consistent;
        server backend1.example.com:12345 weight=5;
        server 127.0.0.1:12345 max_fails=3 fail_timeout=30s;
        server unix:/tmp/backend3;
    }
    upstream dns {
        server 192.168.0.1:53535;
        server dns.example.com:53;
    }
    server {
        listen 12345;
        proxy_connect_timeout 1s;
        proxy_timeout 3s;
        proxy_pass backend;
    }
    server {
        listen 127.0.0.1:53 udp;
        proxy_responses 1;
        proxy_timeout 20s;
        proxy_pass dns;
    }
    server {
        listen [::1]:12345;
        proxy_pass unix:/tmp/stream.socket;
    }
}
```

# stream块ngx\_stream\_core\_module

- nginx stream块的lua扩展
- 目前还处于开发阶段, 还没有正式发布, 需要自己下载安装和体验

文档 <https://github.com/openresty/stream-lua-nginx-module>



# ngx\_stream\_core\_module

预计完成如下特性后会正式发布

- support access\_by\_lua\_block
- access\_by\_lua\_file
- support log\_by\_lua\_block
- log\_by\_lua\_file
- support balancer\_by\_lua\_block
- balancer\_by\_lua\_file
- support ssl\_certificate\_by\_lua\_block
- ssl\_certificate\_by\_lua\_file
- support ngx.semaphore
- add ngx\_meta\_lua\_module
- support lua-resty-core
- add lua\_postpone\_output

# ngx\_stream\_core\_module性能测试

Stream块配置

```
stream {  
    server {  
        listen 1235;  
        content_by_lua_block {  
            local sock = assert(ngx.req.socket(true))  
            local data = sock:receive()  
            ngx.say('hello world')  
        }  
    }  
}
```

```
ab -n 10000 -c 100 -k http://127.0.0.1:1235/
```

# ngx\_stream\_core\_module性能测试

Server Hostname: 127.0.0.1  
Server Port: 1235  
Document Path: /  
Document Length: 0 bytes  
Concurrency Level: 100  
Time taken for tests: 0.406 seconds  
Complete requests: 10000  
Failed requests: 0  
Keep-Alive requests: 0  
Total transferred: 120000 bytes  
HTML transferred: 0 bytes  
Requests per second: 24635.64 [#/sec] (mean)  
Time per request: 4.059 [ms] (mean)  
Time per request: 0.041 [ms] (mean, across all concurrent requests)  
Transfer rate: 288.70 [Kbytes/sec] received

Connection Times (ms)

	min	mean[+/-sd]	median	max
Connect:	0	2 0.7	2	7
Processing:	1	2 0.8	2	9
Waiting:	0	2 0.7	2	8
Total:	2	4 1.1	4	12

# stream lua 的echo 服务器

```
server {  
    listen 1234;  
    content_by_lua_block {  
        local sock = ngx.req.socket()  
        while (true) do  
            local line, err = sock:receive('*l') -- read a line from downstream  
            if not line then  
                ngx.say(err)  
                break  
            elseif line == 'exit' then  
                ngx.say('bye')  
                break  
            end  
            ngx.say(line) -- output data  
        end  
    }  
}
```

# 基于luasocket的客户端

```
local socket = require'socket'
local host = '127.0.0.1'
local port = 1234
local sock = assert(socket.connect(host, port))
sock:settimeout(0)

print('please input anything, exit to quit:')
local input, recvt, sendt, status
local connected = true
```

```
gerrard@ubuntu:~/workspace/test/client$ lua test.lua
please input anything, exit to quit:
hello world
hello world
1234567
1234567
exit
bye
```

```
while connected do
    input = io.read()
    if #input > 0 then
        assert(sock:send(input .. '\n'))
    end
    recvt, sendt, status = socket.select({sock}, nil, 0.1)
    while #recvt > 0 do
        local response, receive_status = sock:receive()
        if receive_status ~= 'closed' then
            if response then
                print(response)
                if response == 'bye' then
                    sock:close()
                    connected = false
                end
            end
        end
    end
    break
end
end
```

# 文档查看工具restydoc

- 可以按模块名搜索，也可以按节名搜，或者二者的组合。
- 搜索方式是精确匹配优先，前缀次之，包含最次

```
gerrard@ubuntu:~/workspace/test/src$ restydoc -h
Usage:
  /home/gerrard/workspace/test/openresty/bin/restydoc [options] [module]

Options:
  -h          Print this help.
  -s SECTION  Specify the section name to be searched

For bug reporting instructions, please see:
  <https://openresty.org/en/community.html>

Copyright (C) Yichun Zhang (agentzh). All rights reserved.
```

# use case

## restydoc -s content\_by\_lua\_file

```
ngx_lua-0.10.6(7)      ngx_lua-0.10.6      ngx_lua-0.10.6(7)
```

**content\_by\_lua\_file**  
**syntax:** `content_by_lua_file <path-to-lua-script-file>`

**context:** `location`, `location if`

**phase:** `content`

Equivalent to `content_by_lua`, except that the file specified by "`<path-to-lua-script-file>`" contains the Lua code, or, as from the "v0.5.0rc32" release, the "Lua/LuaJIT bytecode" to be executed.

Nginx variables can be used in the "`<path-to-lua-script-file>`" string to provide flexibility. This however carries some risks and is not ordinarily recommended.

When a relative path like "foo/bar.lua" is given, they will be turned into the absolute path relative to the "server prefix" path determined by the "-p PATH" command-line option while starting the Nginx server.

When the Lua code cache is turned on (by default), the user code is loaded once at the first request and cached and the Nginx config must be reloaded each time the Lua source file is modified. The Lua code cache can be temporarily disabled during development by switching `lua_code_cache "off"` in "nginx.conf" to avoid reloading Nginx.

Nginx variables are supported in the file path for dynamic dispatch, for example:

```
# WARNING: contents in nginx var must be carefully filtered,
# otherwise there'll be great security risk!
location ~ ^/app/([-_a-zA-Z0-9/]+) {
    set $path $1;
    content_by_lua_file /path/to/lua/app/root/$path.lua;
}
```

But be very careful about malicious user inputs and always carefully

## restydoc resty.redis

```
lua-resty-redis-0.25(7)  lua-resty-redis-0.25  lua-resty-redis-0.25(7)
```

**Name**  
lua-resty-redis - Lua redis client driver for the ngx\_lua based on the cosocket API

This library is considered production ready.

**Description**  
This Lua library is a Redis client driver for the ngx\_lua nginx module:  
<https://github.com/openresty/lua-nginx-module/#readme>

This Lua library takes advantage of ngx\_lua's cosocket API, which ensures 100% nonblocking behavior.

Note that at least ngx\_lua 0.5.14 <https://github.com/chaoslawful/lua-nginx-module/tags> or OpenResty 1.2.1.14 <http://openresty.org/#Download> is required.

**Synopsis**

```
# you do not need the following line if you are using
# the OpenResty bundle:
lua_package_path "/path/to/lua-resty-redis/lib/*.lua;;";

server {
    location /test {
        content_by_lua '
            local redis = require "resty.redis"
            local red = redis:new()

            red:set_timeout(1000) -- 1 sec

            -- or connect to a unix domain socket file listened
            -- by a redis server:
            -- local ok, err = red:connect("unix:/path/to/redis.sock")

            local ok, err = red:connect("127.0.0.1", 6379)
```

# 命令行工具resty

- 像lua和luajit一样的方式运行openresty的lua脚本
- 通过init\_worker\_by\_lua初始化lua代码然后在ngx.timer回调中运行

文档 <https://github.com/openresty/resty-cli>



# usage

```
gerrard@ubuntu:~/workspace/test$ resty -h
resty [options] [lua-file [args]]

Options:
  -c num           Set maximal connection count (default: 64).
  -e prog          Run the inlined Lua code in "prog".
  --help           Print this help.

  --http-include path Include the specified file in the nginx http configuration block
                  (multiple instances are supported).

  -I dir           Add dir to the search paths for Lua libraries.

  --main-include path Include the specified file in the nginx main configuration block
                  (multiple instances are supported).

  --nginx          Specify the nginx path (this option might be removed in the future).
  -V              Print version numbers and nginx configurations.
  --valgrind       Use valgrind to run nginx
  --valgrind-opts  Pass extra options to valgrind

For bug reporting instructions, please see:
  <https://openresty.org/en/community.html>

Copyright (C) Yichun Zhang (agentzh). All rights reserved.
```

# 参考

- 官网 <https://openresty.org/>
- Github <https://github.com/openresty>
- 邮件列表(推荐, 需翻墙)  
<https://groups.google.com/forum/#!forum/openresty>



**The End  
Thanks!**